

Extract, then decide: a modular pipeline of LLMs and rules to identify seizure frequency in epilepsy clinic letters

A dissertation submitted in partial fulfilment of
The requirement for the degree of
MASTER OF SCIENCE in Artificial Intelligence
in
The Queen's University of Belfast
by
Conor Brown
4th September 2026

Declaration of Academic Integrity

I certify that the submission is my own work, all sources are correctly attributed, and the contribution of any AI technologies is fully acknowledged.

I declare that I have read both the University and the School of Electronics, Electrical Engineering and Computer Science guidelines on plagiarism – <https://www.qub.ac.uk/directorates/sgc/learning/LearningResources/Plagiarism/> – and that the attached submission is my own original work. No part of it has been submitted for any other assignment and I have acknowledged in my notes and bibliography all written and electronic sources used.

I am aware of the disciplinary consequences of failing to abide and follow the School and Queen's University Regulations on Plagiarism.

Name : CONOR BROWN

Student Number: 40466218

signature.png

Student's signature:

Date of submission: 4th September 2026

Extract, then decide: a modular pipeline of LLMs and rules to identify seizure frequency in epilepsy clinic letters

Abstract—Seizure frequency is written in clinic letters, and extracting it for research still relies on slow, error-prone manual chart review. Large language models (LLMs) can read these letters, but the best reported result required fine-tuning on over a thousand labelled letters. This paper splits the task in two: an LLM **extracts** every candidate frequency with its evidence, then a fixed policy **decides** the answer. The policy was applied by rules (**Hybrid**) and by a second LLM call (**LLM-only**). Both were evaluated on 1,200 synthetic epilepsy letters against the benchmark reported for the fine-tuned LLM on a different held-out sample of the same corpus. **Hybrid** (F1 0.86) and **LLM-only** (0.85) exceeded the reported benchmark (0.81) and improved on the extraction stage’s answer alone (0.79); removing the extraction prompt’s examples, closed label forms or evidence obligation lowered the result. On this synthetic benchmark, a prompted model with a separate decision step matched a fine-tuned one.

Index Terms—information extraction, hybrid pipeline, large language models, seizure frequency, epilepsy

I. INTRODUCTION

An epilepsy clinic letter is a short narrative, about 400 words, that a specialist writes after seeing a patient. It records the patient’s history, how they are doing now and what treatment is planned. For research, the most useful fact in that letter is usually how often the patient is having seizures: seizure frequency is the main measure of how well epilepsy is controlled and the usual outcome in clinical trials [?]. Epilepsy affects about 1% of the UK population [?], so these letters exist in large numbers, but the frequency is rarely recorded anywhere else in a structured form. A study that needs it must have experts read every letter by hand, which is slow and inconsistent.

Large language models (LLMs) have made this task feasible. Given instructions and a few examples, they can read a letter and report a frequency whether it is written as *2 to 3 seizures a week*, *frequent bouts* or *4 since her last visit* [?], [?]. Two problems remain. First, an LLM’s answer is not always based on the letter in front of it: in one study a model asked for the seizure frequency instead gave advice on the patient’s medication [?]. Second, the most accurate published result on this task came from a model fine-tuned on more than a thousand labelled letters [?], and labelled clinical text is expensive to produce. A method that reached the same accuracy from instructions alone, and showed where in the letter each answer came from, would remove both costs.

There is a further difficulty that is easy to miss. A letter often mentions more than one frequency: the rate last year,

the rate since a medication change, or separate rates for different seizure types. Which of these is the patient’s current frequency is a clinical judgement, and different studies may want different answers. Clinical information extraction has long been described as two steps, finding a mention in the text and converting it to a standard form [?], [?]; this task has a third, choosing between the mentions found. When a single prompt is asked for one final label, all three steps happen inside the model and the choice cannot be seen, checked or changed. This paper separates them. An LLM first **extracts** every frequency mentioned in the letter, together with the text it came from; a fixed, written policy then **decides** which one is the answer. Because the policy works on the extracted record rather than on the letter, it can be applied by rules or by a second LLM call.

The aim of this paper is to find out whether that design can match the fine-tuned benchmark. It asks four questions: whether the pipeline reaches the F1 that Gan *et al.* reported; whether the decision step improves on the answer the extraction call gives on its own; whether it matters if rules or a second LLM call apply the policy; and which parts of the extraction prompt the result depends on. The contribution is the two-stage design, its comparison with the benchmark, and an account of which component the result depends on. The evaluation uses the 1,200 synthetic letters released by Gan *et al.* (Section ??) and their stricter scoring setting (Section ??). On a held-out set of 450 letters, both versions of the pipeline exceeded the reported benchmark (F1 0.86 and 0.85 against 0.81) and both improved on the extraction call’s own answer (0.79). The two studies used different held-out samples from the same corpus, so this is a comparison with a published result rather than a paired test (Section ??), and because the letters are synthetic the results say nothing yet about real clinical text.

The literature review places this split against earlier extractors. The methods describe the letters, the two stages and how they are scored. The results report the benchmark comparison, the lift from decide, the prompt ablations and a check across six models. The discussion and ethical section set out what follows from those results, and what does not.

II. LITERATURE REVIEW

The best approach to extract seizure frequency information from epilepsy clinic letters is an active area of research, with experiments using rules-based methods [?], [?], transformers-based methods [?], [?], [?], [?] and a combination of both [?],

[?] all published this decade. Rules-based methods struggled with the lexical variety of seizure frequency patterns. A key challenge with both rules-based methods and BERT-based methods was their failure to generalise to letters from a different hospital. [?], [?]. LLMs were the first method to demonstrate strong generalisation to a new dataset [?], but this was achieved by fine-tuning a model on over 1,000 synthetic letters. This paper explores whether similar performance can be achieved with a prompted model and rules.

Rules were used to normalise frequencies and dates as a separate layer in two BERT-based pipelines [?], [?], and as minimal repair of malformed LLM output [?], but no study has given rules a decision role after a modern LLM. This paper explores whether LLM instructions and examples combined with decision rules can match a fine-tuned model.

Clinical IE has been defined as the automatic extraction and encoding of clinical information from text [?]: recognising a mention and normalising it into a standard form [?]. Modern LLMs do both in one step when asked for a normalised label such as *4 per month* from a raw statement such as *about once a week*. A further step exists in many clinical extraction pipelines but goes unmentioned: deciding between candidates. In a clinic letter with several seizure frequency statements, producing a single label for the patient’s current status requires the system both to extract the relevant statements and to decide which one best answers the question. Prior rules-based systems made this a distinct stage [?], and clinical concept normalisation routinely separates a high-recall candidate generation stage from a precise ranking stage that chooses among the candidates [?]. LLM pipelines instead make the decision inside the extraction prompt. This paper explores whether separating *extract* from *decide* in a modular pipeline can deliver comparable performance while preserving more of the evidence for downstream use.

Academic papers tend to focus exclusively on benchmark accuracy while practitioners tend to value transparency, configurability and efficiency at the cost of some accuracy [?], [?]. Model interpretability is also a critical factor clinical settings [?]. Clinicians are trained to use deductive logic to draw conclusions and they expect systems to expose transparent, interpretable evidence that they can reliably use to make decisions that demand high accountability. Past work utilises evidence spans as a training input [?] for fine-tuning or as data exhaust as part of using a rule-based method [?]. This paper explores whether preserving a transparent evidence chain from exact text to final answer can make the output both more interpretable and more accurate.

III. METHODS

A. Data

The dataset is a collection of epilepsy clinic letters which record a UK patient’s consultation with an epilepsy specialist, often a neurologist or specialist nurse [?] The NHS requires that clinic letters are sent to the patient’s GP after all outpatient visits such as these. They are often the only record of a patient’s diagnoses, treatments and outcomes in neurological



Fig. 1: An example clinic letter with two distinct seizure frequency patterns highlighted as potential candidates.

TABLE I: Dataset summary

		Development	Test
Number of letters		750	450
Seizure frequency	Frequent	51%	50%
	Infrequent	17%	18%
	Unknown	17%	17%
	Seizure free	15%	15%

disorders, and they are written in narrative form. An example clinic letter is shown in Figure ??.

Gan *et al.* [?] produced this dataset of synthetic letters modelled on real clinic letters from King’s College London hospital. The real letters were used to create 10 template letters and they were combined with 500 short descriptions of seizure frequency patterns produced by GPT-5 to create the full set of 15,099 letters. A sample of 1,200 letters of c. 400 words was used in this paper. Trained experts annotated each letter with a single gold-standard label representing the patient’s current seizure frequency status, and with one supporting quote.

Minimal preprocessing was applied: letters were cleaned only to parse quotes, new lines, tabs, the $\leq$ sign and nulls, and were ingested in full with no tokenisation, section-splitting or truncation.

The dataset was split into a development set of 750 letters and a test set of 450 letters, stratified by their seizure frequency category (frequent, infrequent, unknown, seizure free) to balance the distribution across both splits. The exact breakdown can be seen in Table ??.

TABLE II: What each stage receives and returns

	Extract	Decide
Input	Full letter	Candidate record and decision policy; not the letter
Output	Candidate record: events, each with quoted span, category and canonical label; provisional answer	Final label and its supporting event(s)
Performed by	LLM call 1	Hybrid: rules. LLM-only: LLM call 2

Hybrid and LLM-only share the first call and the extraction record. Full prompts and schemas are in the supporting materials.

B. Architecture

The pipeline has two stages: *extract*, then *decide*. *Extract* is one LLM call that reads the full letter and returns a structured record of every candidate seizure frequency event, each with its exact quoted evidence span, a category and a label in one of the canonical forms used by the gold standard. The same call also proposes a provisional answer from its own candidates. Scoring that answer gives a baseline for a single prompt; the *decide* stage then confirms or replaces it. *Decide* receives that record, not the letter, and produces the final answer by applying a fixed decision policy, either picking one candidate or combining several. Figure ?? illustrates the pipeline with a development-set example.

The decision was performed in two ways on the same extraction record (Table ??). In *Hybrid*, rules apply the policy. In *LLM-only*, a second LLM call applies the same policy, expressed as written instructions and worked examples. The only difference between the two is therefore who performs *decide*; the extraction call, its output and the policy are shared. Because *decide* works from that record, both executors were replayed on the same extraction output without a new model call, so any difference between them is attributable to who performed *decide*.

The pipeline is a fixed sequence of calls and functions. Each LLM call is a prompt with a fixed input and a schema-validated JSON output. A rules-only baseline and several further decompositions of the LLM work were also built and are described in the supporting materials; they are not part of the main comparison.

C. LLM Prompt Design

The most consequential component of the pipeline is the *extract* prompt used for the first call. It has four ingredients, each of which is removed in turn in Section ??:

- 1) *Instructions* to identify relevant seizure frequency events and to propose a provisional answer.
- 2) *Examples* of how events should be represented.
- 3) *Allowed labels*: a closed list of the canonical label forms used by the gold standard.
- 4) An *evidence obligation*: schema fields for the evidence span of every event, with the instruction that each span must be an exact quote from the letter.

TABLE III: Decision policy

When	Do
Several current seizure types are present	Choose the highest current or recent frequency.
The letter gives an overall current count and a breakdown by type	Choose the overall count.
A seizure-free statement sits with other current seizure-like events	Keep seizure-free separate from unknown or last-event statements.

The output schema represents multiple seizure frequency patterns with exact evidence, so that rules or a second LLM call can work from the candidates rather than from the letter. The same decision policy (Table ??) is applied twice: provisionally by the model inside the extraction prompt, and finally at *decide*. For example, if the model finds *3 seizures in March, 6 in May* and provisionally answers with the most recent count, the *decide* stage can replace that with the total over the whole period, *9 per 3 months*, using the evidence the model already found. The second-call *decide* prompt states the same policy as written instructions with worked examples and receives only the candidate record.

The model is instructed to assign each event to one of six categories (a countable or range *frequency rate*, a *cluster frequency*, a *seizure-free* statement, a dated *last event* without a rate, an *unknown frequency* where seizures are discussed without a usable rate, and *no reference*) and to write its label in one of the canonical forms listed in the prompt. The category definitions, label forms and full prompt are in the supporting materials.

D. Rule Design

In *Hybrid*, rules perform *decide*. Each rule is a named function with *regex* matching and if-then logic that reads the candidate record the model returned, not the letter. Rules first rewrite any label the model left in a non-canonical form (a step that changes the score by at most 0.01 and is detailed in the supporting materials), then apply the decision policy in Table ?? . Three kinds of decision rule are used. A *gate* rejects a provisional answer that the policy forbids: if the model finds both *seizure free for eight months* and *two seizures in the past fortnight* but provisionally answers seizure-free, the gate rejects seizure-free while other current seizure-like events remain. A *reselect* rule chooses a different already-extracted event, such as a frequency candidate when the model chose *unknown*. A *rewrite* rule derives a new label from the extracted events: for seizure diaries, date and numeric logic totals the events over the stated period, so *May x0, Jun x0, Jul x1, Aug x0, Sep x1, Oct x2* becomes *4 per 6 months*. Every rule is recorded with the event it acted on, so a final label can be traced to the model’s quoted evidence and the rule that decided it. Rules do not scan the letter for new candidates; every rule acts on the events the model extracted.

E. Pipeline Optimisation

The development set was used to optimise the two key elements of the pipeline: the LLM prompt instructions and the

rules. The focus of the optimisation process was to improve accuracy while keeping the instructions and rules simple and generalisable. Qualitative error analysis was conducted to diagnose the cause of poor performance and identify common error modes. Hypotheses were developed and tested to tune the prompt and rules iteratively. Improving a rule or instruction for one class of error often harmed another, because the annotation policies interact. Separating `extract` from `decide` kept the policy on which facts to extract apart from the policy on which to choose, and made it possible to score the pipeline after each stage and locate bottlenecks. The supporting materials give a fuller account of this process.

F. Evaluation

Every paper reporting seizure frequency extraction this decade used F1 as its primary metric, but they differ in what counts as a correct frequency label. It is not obvious whether *2 per month*, *3 per month* and *once a week* should all be accepted when the gold-standard label is *multiple per month*. To allow comparison with Gan *et al.* [?], this paper adopts their two evaluation settings: a fine-grained `purist` setting and a coarser `pragmatic` setting. If the gold label is *twice a week* and the prediction is *once a week*, that is a correct `pragmatic` answer, since both fall in the `frequent` category, but an incorrect `purist` answer. Table ?? lists the categories in each setting.

Micro-F1 is reported for both settings, because the categories are badly imbalanced (some have over 100 letters, some fewer than 10; Table ??) and micro-F1 gives each letter equal weight. Each letter has exactly one gold label and one prediction, so micro-F1 here equals the share of letters answered correctly. The more precise `purist` F1 is the primary metric.

The pipeline is evaluated on a held-out test set of 450 letters that were never used to tune prompts, rules or any other aspect of the pipeline. Results are compared against the evaluation set labelled Synthetic (1,166) in Gan *et al.* [?], a different held-out sample from the same synthetic corpus, distinct from their fine-tuning data.

The primary model is Gemini 3.7 Flash. It is the only model used for the second decision call (`LLM-only`), for the prompt ablations and for the temperature and thinking-level ablations. Six models were evaluated as the extraction call in `Hybrid` (Section ??). Elsewhere, `Hybrid` results are from Gemini at temperature 0 with low thinking effort. The held-out results are reported as aggregate totals only; individual test letters and their outputs were not inspected.

IV. RESULTS

A. Benchmark Performance

The primary objective of this paper is to evaluate whether a prompted LLM and rules can match a fine-tuned LLM at extracting seizure frequency patterns from epilepsy letters. Results are compared with the Synthetic (1,166) evaluation reported by Gan *et al.* [?]. The two evaluations use identical metric definitions and letters drawn from the same synthetic

TABLE IV: Comparison with previously reported results

Author	Year	Method	Purist F1	Pragmatic
Holgate et al.	2024	Prompted LLM	-	0.73
Holgate et al.	2025	Fine-tuned LLM	0.68	0.81
Gan et al.	2026	Fine-tuned LLM	0.81	0.85
This paper	2026	Hybrid (LLM + rules)	0.86	0.88

Each row was evaluated on a different held-out sample; only Gan et al. and this paper share a corpus and metric definitions.

TABLE V: Hybrid: Purist and Pragmatic categories (test set, $n = 450$)

Purist	n	Purist F1	Pragmatic F1
Daily	41	0.92	Frequent: 0.95
More than weekly, less than daily	123	0.91	
Once a week	5	0.89	
More than monthly, less than weekly	58	0.90	
Once a month	18	0.94	Infrequent: 0.82
More than every 6 months, less than monthly	55	0.81	
Once every 6 months	1	0.00	
Less than once every 6 months	6	0.60	
Unknown	76	0.74	Unknown: 0.74
Seizure free	67	0.90	Seizure free: 0.90

Rows with $n < 10$ shown in grey.

corpus, but different held-out samples. The comparison is therefore with a previously reported benchmark, not a paired evaluation on the same letters, and it does not establish a state-of-the-art result.

Hybrid achieves a `purist` F1 of 0.86 (387/450), above the reported benchmark of 0.81 for the fine-tuned LLM. `LLM-only` also exceeds it (0.85; 383/450). Because micro-F1 equals per-letter accuracy here, an approximate two-proportion comparison gives a difference of +0.05 for Hybrid (95% CI +0.01 to +0.09, $p = 0.01$) and +0.04 for `LLM-only` (95% CI 0.00 to +0.08, $p = 0.04$). These intervals assume independent samples from the same distribution and are indicative rather than a controlled comparison. The first LLM call alone, scored on its provisional answer, reaches 0.79, close to the benchmark. Rules applied to its candidates add +0.07 in Hybrid; the second LLM call adds +0.06 in `LLM-only`.

The pipeline is particularly effective at extracting high-frequency events. Gan *et al.* report `unknown` as their highest scoring category (0.89), with `daily` (0.70) and `more than weekly` (0.68) much lower. In this pipeline, `daily` (0.92) and `more than weekly` (0.91) are among the best performing categories, well above `unknown` (0.74). The full breakdown across `purist` and `pragmatic` categories is in Table ??.

TABLE VI: Purist F1 after each stage, by decision executor

Decision executor	Provisional answer (after extract)	Final answer (after decide)
Hybrid: rules	0.79 (355)	0.86 (387)
LLM-only: second LLM call	0.79 (355)	0.85 (383)

Correct letters of 450 in parentheses. Both rows share the same extraction call and record; only *decide* differs.

B. Provisional and Final Answers

Table ?? reports *purist* F1 at two stops: the provisional answer after *extract*, and the final answer after *decide*, for each decision executor. Each stop scores the answer the pipeline would submit if it stopped there. The provisional answer is the extraction call’s own proposal, so its score measures the quality of that proposal, not candidate identification alone. The difference between the columns is what *decide* adds to the answer already proposed.

The provisional answer from the single extraction call is 0.79. This call is equivalent to the entire pipeline in most prior LLM studies, which ask the model to extract and decide in one prompt. Both executors then improve on that shared provisional answer at *decide*: rules add +0.07 and the second LLM call adds +0.06. This can mean choosing a different candidate, such as a frequency candidate when the model chose *unknown*, or calculating a new label from the chosen evidence, such as totalling the events in a monthly seizure diary. The two executors differ on 28 of 450 letters (rules correct on 16, the second call on 12); a paired McNemar test gives a difference of +0.01 (95% CI −0.01 to +0.03, $p = 0.57$), compatible with no difference. Applying the decision policy to the candidates already extracted is where the pipeline adds value over a single prompt, and on this task it can be performed as well by rules as by a second LLM call. Other decompositions of the LLM work, including a separate label-rewriting call before *decide*, did not improve on this design and are reported in the supporting materials.

C. Extraction Prompt Ablations

Table ?? removes one ingredient of the extraction prompt at a time and scores the *Hybrid* pipeline at both stops, so the effect of each ingredient can be seen on the provisional answer and after rules have applied the policy. Removing the worked examples, the closed list of allowed label forms, or the evidence obligation each lowers the provisional answer by 8 to 10 letters and the final answer by 17 to 23 letters. The evidence-obligation ablation removes the evidence fields and the exact-quote instruction together as one package; it does not test the quote instruction alone. The decrements are not additive and should not be summed: each row removes one ingredient from the full prompt while the others remain. Rules recover some of each loss but never the whole of it, which shows that the decision stage depends on the quality of the record the extraction call produces.

TABLE VII: Extraction prompt ablations (Hybrid, Gemini, test set)

Extraction prompt	Provisional answer	Final answer
Full prompt	0.79 (355)	0.86 (387)
Without examples	0.77 (345)	0.82 (370)
Without closed allowed-label forms	0.77 (347)	0.81 (364)
Without evidence obligation	0.77 (345)	0.82 (370)

Purist F1 with correct letters of 450 in parentheses. Each row changes only the extraction prompt; the same rules apply the policy at *decide*. Effects are not additive.

D. Model Comparison

Hybrid shows consistent, incremental gains on each of the six models evaluated as the extraction call (Figure ??). The average gain from the provisional answer to the final answer was +0.07, with *Luna* showing the largest gain of +0.10. *Purist* F1 closely tracks general medical capability as measured by Artificial Analysis’ Healthcare and Medical Index. The models cluster into three groups: strong (Gemini 0.86, Grok 0.85), medium (DeepSeek 0.82, *Luna* 0.79) and weak (Qwen 0.76, Gemma 0.72), which largely reflect model size and cost.

Table ?? reports how well each model kept to the extraction contract. Every model produced a parseable record for at least 442 of 450 letters, and no call failed. The four hosted models quoted an exact substring of the letter as evidence for the selected event in 99% or more of letters; the two local models paraphrased more often (Qwen 97%, Gemma 93%) and needed more non-semantic format repair. An exact span is not proof that the right statement was chosen, but it means that a reviewer can locate the evidence for almost every answer verbatim in the source.

Table ?? also lists the two earlier papers that used the same *pragmatic* metric on different data. Holgate *et al.* [?] ran Llama 2 on local hospital servers with eight Nvidia V100 GPUs and reached a *pragmatic* F1 of 0.73 without fine-tuning. In *Hybrid*, Qwen reached a *pragmatic* F1 of 0.80 running locally on a laptop with a single GPU. This shows that the design can execute on local hardware under the same synthetic task conditions, not that it performs on real letters or is fit for deployment.

E. Model Configuration

Raising the reasoning effort level did not improve performance. Gemini’s *purist*-F1 score was 0.86 at low and 0.84 at high; a paired difference of +0.02, 95% CI −0.01 to 0.04; $p = 0.31$. Extra reasoning adds cost and latency without improving extraction.

The temperature setting also had no significant effect: 0.86 at temperature 0 versus 0.84 at temperature 1 (paired difference +0.02, 95% CI −0.01 to 0.04; $p = 0.20$). Temperature 0 is therefore an appropriate default, and *Luna*’s fixed temperature of 1 does not invalidate its comparison. The task is tightly defined by an explicit schema, which limits the expected

TABLE VIII: Six models as the extraction call: scores and contract adherence (Hybrid, test set, $n = 450$)

Model	Purist	Prag.	Exact ev.	Repairs	Unparsed
Gemini 3.7 Flash	0.86	0.88	99.8%	0	1
Grok 4.6	0.85	0.89	100%	0	0
DeepSeek V4 Flash	0.82	0.85	99.6%	2	0
GPT-5.6 Luna	0.79	0.82	98.9%	0	1
Qwen 3.8 27B (local)	0.76	0.81	96.7%	3	2
Gemma 4 26B (local)	0.72	0.77	92.9%	15	8

Exact ev.: the quoted span for the selected event is an exact substring of the letter. Repairs: letters needing a non-semantic JSON dialect repair. Unparsed: letters with no valid record after repair, scored as incorrect. No call failed for any model.

variability by design and partly explains why neither setting mattered.

F. Error Analysis

The worst performing category was unknown with an F1 score of 0.74. 69% of the errors (37/54) are incorrect answers of unknown. The full error distribution is shown in the confusion matrix (Figure ??).

The three most common errors are infrequent $\rightarrow$ unknown (16), frequent $\rightarrow$ unknown (13), and seizure free $\rightarrow$ unknown (8). The errors reflect a cautious approach when the evidence is ambiguous, for example the model recognised the phrase ‘these episodes crop up most shifts’ as the primary seizure frequency description but refused to make an assumption about how many shifts the person works. The gold-standard label was multiple per week but the model chose unknown. This cautious assessment is justifiable, and because the exact evidence is preserved, downstream users can see the proposed answer and draw their own conclusion from the quoted span.

G. Experimental Environment

Table ?? summarises the hardware, software and model settings used. DSPy co-ordinated the model calls and required structured JSON output, Pydantic validated each record against the schema, and LiteLLM delivered API calls to Ollama and the model routers. Full specifications, API settings and replay artefacts are in the supporting materials.

V. DISCUSSION

This paper shows that, on the Gan synthetic benchmark, one prompted LLM call and a decision policy extract an epilepsy patient’s current seizure frequency pattern at least as accurately as the fine-tuned LLM previously reported on a different held-out sample of the same corpus [?]. A single prompt designed to extract every candidate with exact evidence, followed by rules or a second call that decide from that evidence, matched that benchmark. The same design also executed on a laptop-class GPU with open-weight models, at a lower but still competitive score. Later researchers or healthcare organisations can

TABLE IX: Experimental environment

Component	Setting
Local models	Qwen 3.8 27B and Gemma 4 26B served by Ollama 0.32.15 on a Dell XPS 16 9640 (Windows 11, Intel Core Ultra 9 185H, 32 GB RAM, NVIDIA RTX 4070 Laptop GPU, 8 GB VRAM)
Hosted models	Gemini 3.7 Flash, Grok 4.6, GPT-5.6 Luna and DeepSeek V4 Flash, called from a Mac mini M1 (16 GB, macOS 26.5.2) directly (OpenAI, DeepSeek) or via OpenRouter and AI Gateway
Model settings	Temperature 0 (Luna accepts only 1); thinking or effort level low. Gemini ablations: temperature 1; thinking medium and high
Software	Python 3.14.4, DSPy 3.2.1, Pydantic 2.13.4, LiteLLM 1.87.1

evaluate it on their own data and hardware; it is not ready for local deployment.

The method is particularly effective at identifying the categories that mark the greatest disease burden. F1 scores of 0.92 for daily and 0.91 for more than weekly are around 0.2 higher than the category scores previously reported [?], although on a different sample. Structured seizure frequency of this kind is intended for later research uses: cohort identification for trials, retrospective and longitudinal analyses of treatment response, and outcome variables for predictive models [?], which together can widen the scope of real-world evidence [?]. Those are intended uses, not capabilities demonstrated here; each requires validation on real records first.

Past work has typically tasked the LLM with producing a single answer in a fully normalised form [?]. That discards the semantic depth of the narrative, which clinicians value because it conveys reasoning, uncertainty and vague impressions, and those nuanced observations are more accurate and reliable [?]. Current LLMs can populate a richer record that preserves them: every candidate with its quoted span, the policy applied, and the label it produced. The ablations show that this record is also what the accuracy depends on. A more explainable system is often a more useful system in a clinical setting [?], and here it is not a less accurate one.

There are several limitations. Synthetic data was used for all evaluations, and while past research has shown models fine-tuned on this dataset generalise effectively to real patient data [?], the methods tested here remain untested on real letters; no claim is made about clinical validity, workflow fit or privacy compliance. The held-out samples differ. Past research has shown that rules tuned on one corpus can generalise poorly to others, so the gains from rules may not transfer to other datasets [?]. The stage scores are cumulative stops rather than isolated stage metrics. Evidence quality was measured only as exact-substring adherence, which does not show that a span semantically supports the label; the gold standard’s single reference quote is a loose comparator for that purpose, since it may be paraphrased or abbreviated. A development-set study of semantic evidence sufficiency, adjudicated from the extracted spans and the policy alone, is specified as the next step before any further held-out evaluation. Finally, the

method was applied to one narrow task. The same method has been configured for a multi-fact epilepsy inventory task [?], to be reported separately, and future work should evaluate whether the design generalises to other clinical IE tasks and whether a more general framework, in which an LLM chooses how to represent facts and which policy to apply, adds value as reported in oncology [?] and cardiology [?].

VI. AI ETHICAL CONSIDERATIONS

This evaluation used only the synthetic clinic letters released by Gan *et al.* [?]. Those letters are modelled on real UK epilepsy letters but do not contain identifiable patient records. Hospitals prefer processing patient data on their own servers as it is highly sensitive [?]. Two of the six models were run entirely on local hardware to test whether the design can execute in such a setting on synthetic data, not as a privacy-compliance assessment.

LLMs can invent evidence, so careful consideration was given to the prompt instructions and schema design. The model is required to quote exact evidence spans from the letter, and every later change to the answer is made by a named rule or a second call applied to that quoted evidence. A reviewer can therefore trace any final label back to the text that supports it and to the policy that decided it [?]. This transparency matters more than headline accuracy in a clinical setting, where an unexplained answer cannot be safely acted on.

The letters contain no demographic attributes, so unequal performance across patient groups cannot be tested, and the synthetic corpus may not reflect the language of every clinic. Bias can enter through the models, the prompts and the decision rules; for example, the cautious default to `unknown` (Section ??) would systematically under-count patients whose seizures are described vaguely. Rules make the decision policy inspectable and changeable by local clinicians, but they do not remove the need for local clinical review and fairness testing on real records before any operational use.

VII. CONCLUSION

Seizure frequency is a core outcome measure in epilepsy, recorded almost exclusively in narrative clinic letters. This paper proposes a pipeline in which one LLM call extracts every candidate fact with its quoted evidence in a canonical form, then a decision applies a fixed policy to those candidates.


Both Hybrid (Purist F1: 0.86) and LLM-only (0.85) exceeded the benchmark previously reported for a fine-tuned LLM (0.81) on a different held-out sample of the same synthetic corpus. The provisional answer reaches 0.79; applying the decision policy to the evidence it found adds a further +0.06 to +0.07, whether that policy is applied by rules or by a second LLM call. The result depends on the extraction prompt's examples, closed label forms and evidence obligation, and it held across six models, including two run on a single laptop GPU.

Automatically extracting seizure frequency patterns from electronic health records could make this outcome measure available for large-scale clinical research: cohort identification,

retrospective studies and predictive modelling. Future research should evaluate it on real patient data and on a broader set of clinical information extraction tasks before any of those uses.

REFERENCES

- [1] S. Wigglesworth, A. Neligan, J. M. Dickson, A. Pullen, E. Yelland, T. Anjuman, and M. Reuber, "The incidence and prevalence of epilepsy in the United Kingdom 2013–2018: A retrospective cohort study of UK primary care data," *Seizure*, vol. 105, pp. 37–42, 2023.
- [2] R. S. Fisher *et al.*, "ILAE official report: A practical clinical definition of epilepsy," *Epilepsia*, vol. 55, no. 4, pp. 475–482, 2014.
- [3] Y. Gan *et al.*, "Reproducible synthetic clinical letters for seizure frequency information extraction," arXiv:2603.11407, 2026.
- [4] B. Fonferko-Shadrach *et al.*, "Annotation of epilepsy clinic letters for natural language processing," *J. Biomed. Semantics*, vol. 15, no. 17, 2024.
- [5] B. Holgate, J. Davies, S. Fang, J. S. Winston, J. T. Teo, and M. P. Richardson, "Fine-tuning LLMs to extract epilepsy seizure frequency data from health records," in *Proc. 24th Workshop on Biomedical Language Processing (BioNLP)*, 2025, pp. 44–55.
- [6] B. M. Decker *et al.*, "Development of a natural language processing algorithm to extract seizure types and frequencies from the electronic health record," *Seizure*, vol. 101, pp. 48–51, 2022.
- [7] K. Xie *et al.*, "Extracting seizure frequency from epilepsy clinic notes: a machine reading approach to natural language processing," *J. Amer. Med. Informat. Assoc.*, vol. 29, no. 5, pp. 873–881, 2022.
- [8] B. Holgate, S. Fang, A. Shek, M. McWilliam, P. Viana, J. S. Winston, J. T. Teo, and M. P. Richardson, "Extracting epilepsy patient data with Llama 2," in *Proc. 23rd Workshop on Biomedical Natural Language Processing (BioNLP)*, 2024, pp. 526–535.
- [9] K. Xie, B. Litt, D. Roth, and C. A. Ellis, "Quantifying clinical outcome measures in patients with epilepsy using the electronic health record," in *Proc. BioNLP Workshop*, 2022, pp. 369–375.
- [10] M. Fernandes, A. Cardall, L. M. V. R. Moura, C. McGraw, S. F. Zafar, and M. B. Westover, "Extracting seizure control metrics from clinic notes of patients with epilepsy: A natural language processing approach," *Epilepsy Res.*, vol. 207, 2024, Art. no. 107451.
- [11] T. Brown *et al.*, "Language models are few-shot learners," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 33, 2020, pp. 1877–1901.
- [12] M. Agrawal, S. Hegselmann, H. Lang, Y. Kim, and D. Sontag, "Large language models are few-shot clinical information extractors," in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP)*, 2022, pp. 1998–2022.
- [13] L. Cui, S. S. Sahoo, S. D. Lhatoo, G. Garg, P. Rai, A. Bozorgi, and G.-Q. Zhang, "Complex epilepsy phenotype extraction from narrative clinical discharge summaries," *J. Biomed. Informat.*, vol. 51, pp. 272–279, 2014.
- [14] S. Fu *et al.*, "Clinical concept extraction: A methodology review," *J. Biomed. Informat.*, vol. 109, 2020, Art. no. 103526.
- [15] Y. Wang *et al.*, "Clinical information extraction applications: A literature review," *J. Biomed. Informat.*, vol. 77, pp. 34–49, 2018.
- [16] L. Chiticariu, Y. Li, and F. Reiss, "Rule-based information extraction is dead! Long live rule-based information extraction systems!," in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP)*, 2013, pp. 827–832.
- [17] R. E. Sherman *et al.*, "Real-world evidence—what is it and what can it tell us?" *N. Engl. J. Med.*, vol. 375, no. 23, pp. 2293–2297, 2016.
- [18] S. T. Rosenbloom, J. C. Denny, H. Xu, N. Lorenzi, W. W. Stead, and K. B. Johnson, "Data from clinical notes: A perspective on the tension between structure and flexible documentation," *J. Amer. Med. Informat. Assoc.*, vol. 18, no. 2, pp. 181–186, 2011.
- [19] Z. Yang *et al.*, "CLINES: Clinical LLM-based information extraction and structuring agent," medRxiv:10.64898/2025.12.01.25341355, 2025.
- [20] J. Wang *et al.*, "Scaling clinician-grade feature generation from clinical notes with multi-agent language models," arXiv:2508.01956, 2025.
- [21] D. Xu, M. Gopale, J. Zhang, K. Brown, E. Begoli, and S. Bethard, "Unified Medical Language System resources improve sieve-based generation and Bidirectional Encoder Representations from Transformers (BERT)-based ranking for concept normalization," *J. Amer. Med. Informat. Assoc.*, vol. 27, no. 10, pp. 1510–1519, 2020.




pipeline_architecture.png

Fig. 2: A clinic letter is processed by `extract`, then `decide`, to produce a structured label with grounding evidence. `Extract` is one LLM call. `Decide` works from the resulting record and is performed by rules (Hybrid) or by a second LLM call (LLM-only).



six_model_stage_performance.pdf

Fig. 3: Hybrid Purist F1 for six models: the LLM’s provisional answer after `extract` and the final answer after rule-based `decide`. Every model gains at `decide`; Luna gains most (0.69 to 0.79).



confusion_matrix_pragmatic.pdf

Fig. 4: Hybrid confusion matrix on the pragmatic categories (test set). Most errors are incorrect predictions of unknown.