

Modular and Auditable Clinical Extraction from Epilepsy Letters

Author Names Omitted for Draft

Affiliations Omitted for Draft

City, Country

email@example.com

Abstract—We evaluate a modular clinical extraction architecture on two complementary epilepsy-letter tasks: deep seizure-frequency extraction and broad clinical phenotyping. The first task tests temporal reasoning and benchmark-facing normalization on the Gan 2026 seizure-frequency corpus; the second tests multi-entity recovery on ExECTv2 using de-duplicated clinical-headline facts as the primary surface. Across both tasks, the system separates model-mediated clinical interpretation, deterministic normalization and projection, evidence-grounding checks, reliability audits, and component-impact analyses. This draft reports aggregate validation and frozen holdout evidence under fixed inspection boundaries, and it deliberately separates reliability scorecards from component-causal claims.

Index Terms—clinical information extraction, epilepsy, large language models, hybrid systems, reliability, component ablation

I. INTRODUCTION

Clinical extraction systems must often trade breadth, depth, auditability, and portability. Epilepsy letters make the tension concrete: seizure-frequency statements require temporal anchoring, seizure-free handling, current-versus-historical selection, and benchmark-compatible normalization, while phenotyping requires the same infrastructure to recover diagnosis, medication, investigation, and seizure-frequency facts without collapsing into a task-specific rule stack. Prior epilepsy NLP has shown both the value of broad clinic-letter extraction and the persistent difficulty of seizure-frequency concepts [1], [2], while recent work emphasizes pretrained models, LLM fine-tuning, and synthetic but task-faithful supervision [3]–[6].

This paper studies a modular extraction design in which deterministic rules and large language model (LLM) components are explicit, testable parts of the scientific object, evaluated on a deep single-concept task (Gan 2026 seizure-frequency extraction) and a broad multi-entity task (ExECTv2 clinical-fact recovery). The central claim is not that any one model or rule set is universally optimal, but that explicit component boundaries let us report performance, reliability, and component impact without blending incompatible evidence surfaces.

II. RELATED WORK

Prior epilepsy natural language processing includes broad phenotyping systems and narrow seizure-frequency extraction systems. Recent LLM-based clinical extraction has improved coverage and flexibility, but can obscure the role of deterministic preprocessing, schema repair, normalization, and post-processing [4], [5]. We instead position deterministic rules

as controlled variables rather than incidental implementation detail, and treat reliability audits and component-impact replays as separate manuscript objects. This emphasis aligns with evidence-grounded seizure-frequency supervision in recent synthetic-letter work, extending its reporting discipline to component boundaries and aggregate replay evidence [6].

III. METHODS

A. Dataset Construction and Scoring Surfaces

Gan 2026 is used as the deep seizure-frequency task [6]. The working corpus is the released 1,500-letter synthetic seizure-frequency set, partitioned once into train, validation, and locked test manifests. Train rows are reserved for optimizer-only procedures, validation rows are the development surface, and `test450` is reserved for frozen final aggregate evaluation. Gan outputs are rendered to the task label grammar and scored under the Purist and Pragmatic surfaces. Purist is the stricter final-label surface; Pragmatic is a companion read of acceptable clinical-label equivalence. Test-set row-level failure inspection is not used for development or tuning.

ExECTv2 is used as the broad phenotyping task, building on the ExECT lineage of epilepsy clinic-letter extraction and the later synthetic-letter annotation resource [1], [2]. The ExECTv2 working sets are treated as clinical-letter assemblies with entity-specific annotations for Diagnosis, SeizureFrequency, Prescription, and Investigations. The primary surface is de-duplicated `clinical_headline` recovery: predictions and gold annotations are projected to one clinical-fact view per headline assertion before scoring. Strict benchmark and CUI-oriented outputs are retained as diagnostic compatibility surfaces with the original annotation contract, not as the primary success criterion.

B. Architecture Families

We compare rules-only, LLM-only, and hybrid families. In Gan 2026, the promoted hybrid structured-event architecture uses an LLM to extract source-near seizure events and deterministic code for normalization, projection, rendering, and scoring. In ExECTv2, a manifest-driven clinical-finding assembly combines candidate producers, entity-specific reconciliation lenses, evidence validation, and scoring views.

C. Inspection Boundaries and Claim Classes

Each result row carries an inspection boundary. Gan validation artifacts support development and row-level error analysis; Gan locked-test artifacts are aggregate final-evaluation evidence only unless a separate post-hoc protocol is declared. ExECTv2 dev140 artifacts support row-level development and model-swap diagnostics; full-200 aggregate artifacts authorize only aggregate reliability and component-impact reads, not full-200 row-level failure tuning.

We distinguish four evidence types throughout. Architecture comparison describes aggregate task performance under a fixed scorer. Frozen holdout evidence reports locked evaluation outputs without authorizing post-hoc tuning. Reliability scorecards test whether a fixed system is grounded, calibrated, robust, consistent, and operationally replayable. Component-impact analysis is reserved for ablations or same-input replays that measure a named component’s score contribution under the declared scorer, split, and inspection boundary. Appendix Table VI records the source artifacts and claim boundary for the main result families so that frozen aggregate evidence, development evidence, reliability evidence, and component-impact evidence remain traceable and non-interchangeable.

IV. RESULTS

We evaluate the shared modular architecture on two complementary epilepsy-letter tasks: deep seizure-frequency extraction (Section IV-A) and broad multi-entity phenotyping (Section IV-B). For each task we report architecture comparison, frozen holdout or aggregate validation evidence under declared inspection boundaries, and, where predeclared, separate reliability and component-impact readouts.

A. Gan 2026 Seizure-Frequency Extraction

1) *Three-Way Architecture Comparison:* On validation750, `hybrid_structured_events` combined the best LLM-using coverage and accuracy balance among the canonical families: the LLM extracts source-near structured events with exact evidence spans, and deterministic code owns normalization, projection, and scoring. The deterministic canonical pipeline led validation after de-overfitting but showed the largest validation-to-holdout drop on the frozen `test450` audit, supporting the portability warning that high validation score from rules alone is incomplete generalization evidence. The validation750 table is retained as supporting development evidence in Appendix Table VII.

The promoted close-off candidate is the single GPT structured-event pass on `gpt-4.1-mini`, which reached 364/450 Purist (0.809) on locked `test450` with the smallest validation-to-test drop among LLM-using architectures. The full V12 fresh-evidence hybrid reached 379/450 Purist (0.842) but is retained only as a high-complexity ceiling comparator, not the operational headline system.

2) *Frozen Holdout Aggregate Evidence:* Late-cycle consensus/fresh selector evidence is reported only as frozen aggregate audits completed under predeclared Gate 4 protocols. The

constrained-source audit failed promotion bars; the exact-source audit passed promotion bars. Neither audit authorizes post-test tuning or row-level failure inspection. The audit rows are appendix evidence because they are late-cycle selector diagnostics rather than the promoted operational architecture (Appendix Table VIII).

Gan holdout-facing claims are limited to aggregate `test450` numbers under frozen protocols. Validation750 rows are development evidence. The structured-event hybrid leads the frozen three-way comparison, while the exact-source consensus/fresh selector records a bounded +16 Purist holdout gain at 0.6000 changed-label precision. These results do not authorize further gate, prompt, or selector tuning from locked-test aggregates.

3) *Bridge to ExECTv2:* Seizure-frequency normalization, temporal anchoring, and seizure-free handling developed for Gan 2026 are the direct substrate for ExECTv2 SeizureFrequency recovery. The cross-task read is not a benchmark-win claim: it tests whether the modular investment transfers to a new annotation schema and scoring surface. ExECTv2 SeizureFrequency remains the weakest family on `clinical_headline` despite positive full-200 aggregate recovery, consistent with the deep-reasoning difficulty established by Gan 2026.

B. ExECTv2 Clinical-Fact Recovery

1) *Architecture and Clinical-Headline Performance:* We evaluate the selected ExECTv2 clinical-finding assembly primarily by de-duplicated clinical-headline recovery rather than strict full-schema annotation reproduction. On the aggregate full-200 current-code audit, the verifier-backed GPT-4.1-mini v08-shaped architecture scored 0.8680 overall clinical-headline F1 (0.8502 before the 2026-07-02 scorer-correctness sweep and same-day P7 deterministic-producer fix; see Table II footnote). The accepted lean candidate, the 2-call no-SF-adjudicator GPT-4.1-mini system, reduced the full-200 call profile to 400 calls while preserving the governing cost-performance gates: 0.8356 overall and 0.7525 SeizureFrequency F1 (pre-sweep scoring; not yet re-scored under the finalized scorer).

Strict benchmark and CUI-oriented outputs measure compatibility with the original annotation surface. They are reported only as secondary diagnostics and are not described as the primary ExECTv2 success criterion.

2) *Same-Core Model Swap:* Using the frozen `exectv2_2call_no_sf_adjudicator_model_swap` core, DeepSeek chat, GPT-4.1-mini, and Qwen 3.6 35B were compared under the same component graph and `clinical_headline` scorer. On dev140, DeepSeek produced the strongest aggregate score, GPT-4.1-mini remained operationally clean, and the unrepaired Qwen row remained diagnostic because output-contract failures blocked operational promotion. Qwen repair v02 later passed the predeclared dev140 and full-200 aggregate gates with zero call/parse failures in both assemblies. Detailed model-swap rows move to the appendix because they support model-family




Fig. 1. Two-task modular evidence architecture. Gan 2026 stresses temporal seizure-frequency extraction under frozen aggregate holdout rules; ExECTv2 reuses the same separation of model-mediated candidates, deterministic clinical logic, scoring views, and reliability/component-impact evidence for broad clinical-fact recovery.

TABLE I
GAN 2026 FROZEN TEST450 AGGREGATE AUDIT (GPT-4.1-MINI)

Architecture	Rendered	Purist (of rendered)	Pragmatic (of rendered)	Claim use
deterministic_canonical_pipeline	450/450	329/450 (0.731)	341/450 (0.758)	Frozen aggregate; validation-to-test gap warning.
hybrid	334/450	269/334 (0.805)	281/334 (0.841)	Frozen aggregate; coverage-limited.
hybrid_structured_events	448/450	364/448 (0.812)	381/448 (0.850)	Strongest frozen hybrid aggregate on rendered rows.
llm_only_canonical_pipeline	450/450	326/450 (0.724)	346/450 (0.769)	Frozen comparator row.

TABLE II
EXECTv2 ARCHITECTURE COMPARISON ON FULL-200 AGGREGATE CLINICAL_HEADLINE. THE V08 ROW REFLECTS THE FINALIZED SCORER (2026-07-02 FOUR-FAMILY SCORER-CORRECTNESS SWEEP + SAME-DAY P7 DETERMINISTIC-PRODUCER FIX; OVERALL 0.8502→0.8680). THE NO-VERIFIER AND LEAN 2-CALL ROWS ARE RETAINED AT THEIR PRE-SWEEP VALUES (NOT YET RE-SCORED UNDER THE CURRENT SCORER) AS COMPONENT-ROLE AND COST-FRONTIER COMPARATORS; THEIR RELATIVE ORDERING IS UNCHANGED.

System/view	Overall	Diagnosis	SF	Prescription	Investigations	Claim use
Current-code v08-shaped GPT-4.1-mini	0.8680	0.8546	0.7842	0.9278	0.9213	Aggregate architecture evidence.
No-verifier ablation [†]	0.8431	0.8410	0.7850	0.8926	0.8563	Component-role comparison.
Accepted lean 2-call no-SF-adjudicator [†]	0.8356	0.8397	0.7525	0.8926	0.8563	Cost-performance frontier.

[†]Pre-sweep scoring; not yet re-scored under the finalized scorer.

robustness rather than the central architecture comparison (Appendix Tables IX and X).

The same-core model-swap study suggests that the ExECTv2 component graph is not GPT-4.1-mini-specific. DeepSeek reached 0.8643 and GPT-4.1-mini reached 0.8526 clinical-headline F1 on dev140 under the frozen core (re-scored under the finalized 2026-07-02 scorer), with full-200 aggregate rows of 0.8566 and 0.8356 respectively. Qwen repair v02 provides additional same-core model-family evidence after output-contract repair, but remains below the operational GPT-4.1-mini and DeepSeek rows. The unrepaired Qwen dev140 row is retained as a diagnostic comparator only.

3) *Reliability Scorecard*: The reliability scorecard follows the main performance tables and tests whether the selected ExECTv2 system behaves faithfully under fixed scoring and inspection boundaries. Current evidence supports aggregate full-200 reliability claims with no holdout or deployment-probability claim.

The reliability scorecard tests whether the selected ExECTv2 system behaves faithfully under fixed scoring and inspection boundaries. Calibration and robustness have predeclared aggregate full-200 evidence; review routing has a useful high-recall diagnostic but no validated low-burden policy; and consistency is supported by saved live-repeat panels for the

selected lean candidate. These results support trust in the fixed architecture but are not holdout performance claims and are not component-causal ablation evidence.

4) *Component Impact*: Component impact is reported separately from the reliability scorecard. Claims are limited to aggregate replay deltas under a fixed scorer, split, and inspection boundary. Dev140 one-component-off replay and full-200 aggregate-only replay support limited component-impact language for dictionary normalization, residual semantic recovery, and headline projection on `clinical_headline`.

The no-verifier full-200 ablation and Investigations rule ablation remain separate component-role comparisons. They should not be blended with the layer-ladder component-off readouts above. Component-impact claims are reserved for ablations and same-input replay deltas under a declared scorer and inspection boundary. Dictionary, residual-semantic, and headline-projection layers show positive aggregate deltas on dev140 and full-200 replay, but these results do not prove any component is globally required.

5) *Gap Mechanism: Gold-Quality Ceilings*: A whole-corpus row-level adjudication of every `clinical_headline` disagreement, scored under the finalized 2026-07-02 scorer, separates genuine model error from gold-quality and scorer-mechanics artifacts. The finding

TABLE III
EXECTV2 RELIABILITY SCORECARD

Dimension	Current evidence	Manuscript claim
Evidence grounding	Same-core rows report minimum exact evidence rate 1.0000; robustness audit has schema/evidence validity 1.0000.	Outputs are evidence-grounded on audited surfaces.
Calibration	Frozen grouped scoring rule validates on full-200 with ECE 0.0432, Brier 0.2245, and base-rate Brier 0.2387.	Aggregate calibration evidence.
Review routing	Lower-burden candidate failed validation because burden rose to 0.9661 despite 0.9037 catch.	No promoted low-burden triage policy.
Robustness	Current-code v08 hard-slice F1 is 0.8336 across 414 eligible family cells versus 0.8503 overall.	Aggregate hard-slice validation evidence.
Consistency	Lean candidate has hard50 temp-0 exact agreement 0.9217 and dev140 varying-temperature exact agreement 0.8857.	Live-repeat consistency evidence.
Component role	Investigations verifier plus deterministic suppression remains strongest at 0.9213; deterministic replacement is not ready.	Rules are verification/suppression aids.

TABLE IV
COMPONENT-OFF AGGREGATE DELTAS ON `CLINICAL_HEADLINE`

Component	Type	Split	Overall delta	Main family signal	Claim boundary
<code>standard_dictionary</code>	dictionary	dev140	+0.0389 to +0.1120	Diagnosis/SF gains	Conditional dictionary recovery.
<code>standard_dictionary</code>	dictionary	full200	+0.0186 to +0.0290	Diagnosis up to +0.0802	Same-core full-200 impact only.
<code>residual_semantic_lens</code>	semantic lens	dev140	+0.0175 to +0.1041	Investigations gains	Prediction-bearing semantic contribution.
<code>residual_semantic_lens</code>	semantic lens	full200	+0.0098 to +0.0117	Diagnosis up to +0.0310	Aggregate replay only.
<code>headline_projection</code>	deterministic projection	dev140	+0.0283 to +0.0446	SF up to +0.2031	Projection/format contribution.
<code>headline_projection</code>	deterministic projection	full200	+0.0302 to +0.0350	SF up to +0.1417	Format layer only.

is family-specific and inverts the naive reading of the per-family F1 gap.

Diagnosis carries the most disagreements (199 on dev140), but only **15.6%** (31/199) are genuine model errors; the remaining 84.4% are gold-quality artifacts — dominantly gold multiplicity, where the gold tags both a parent and a child concept (or splits one compound diagnostic phrase into separate atomic tags) from a single statement and the model’s reasonable one-tag consolidation is scored wrong. The D1 hierarchy-aware concept match folded 5 such parent/child pairs into the metric (F1 0.6617→0.6779), reinforcing rather than threatening the gold-artifact finding.¹

SeizureFrequency is the corpus’s lowest-agreement entity (human IAA 0.47). Of its 64 adjudicated disagreements, **29.7%** (19/64) are genuine; the rest are gold under-annotation, redundant double-tagging, and inter-annotator ambiguity. Counting only genuine errors, the program is clinically defensible on far more letters than the raw metric credits. The metric carries ± 0.03 run-to-run variance from temperature-0 nondeterminism, so SF figures are reported as bands.

Prescription is the exception: **72.2%** (26/36) of its disagreements are genuine model errors (polypharmacy drug omissions, current-vs-future dose conflation, non-AED over-extraction) — the opposite pattern, where the model is the bottleneck, not the gold.

¹Diagnosis is reported on the `concept_only` surface (entity-agnostic recall, home-tagged precision) — the surface `clinical_headline_unit_keys` actually maps Diagnosis to, and the surface the production assembly’s `score_for_primary` uses for the overall. The companion `concept_negation` surface (0.6617, negation-aware) is retained as a diagnostic only; the finalized overall uses `concept_only` (0.6779) so the headline equals the micro-average of the four cited per-family surfaces.

TABLE V
DETERMINISTIC VS LLM-TUNED PRESCRIPTION PRODUCER THROUGH THE V08 ASSEMBLY (RX `CLINICAL_HEADLINE`).

Split	Producer	F1	P	R	TP/FP/FN
dev140	Deterministic pre-P7	0.9386	0.9502	0.9272	191/10/15
dev140	LLM-tuned	0.9526	0.9795	0.9272	191/4/15
dev140	Deterministic P7-fixed	0.9615	0.9524	0.9709	200/10/6
full-200	Deterministic pre-P7	0.9033	0.9312	0.8771	257/19/36
full-200	Deterministic P7-fixed	0.9278	0.9343	0.9215	270/19/23
full-200	LLM-tuned	0.9492	0.9748	0.9249	271/7/22

Deterministic vs LLM-tuned Prescription (comparator). The v08 architecture uses a deterministic Prescription lane (AED-only by lexicon construction, with per-clause current-vs-future gating). A best-effort LLM extractor (the canonical evolved instruction plus current-vs-future and AED-only behavior probes, the latter with an emit-if-unsure safety clause) was run through the same assembly as a comparator. The result is split-dependent: on dev140 the deterministic producer wins (0.9615 vs 0.9526, recall-driven); on full-200 the LLM-tuned extractor wins (0.9492 vs 0.9278 Rx F1; 0.8730 vs 0.8680 overall, precision+recall). The two approaches fix different failure modes with different dev/test prevalence — the deterministic P7 fix targets multi-dose weight-context over-suppression (recall); the LLM AED-only gate targets non-AED comorbidity-drug over-extraction (precision) that the deterministic lexicon does not fully handle on the broader test surface. The deterministic lane is competitive with, not strictly dominant over, a tuned LLM on Prescription; the full-200 LLM overall advantage (+0.0050) is within run-to-run variance and is not over-claimed. Table V.

Investigations: 65.7% (23/35) genuine, concentrated in the

MRI-crowds-out-EEG recall omission pattern.

The implication for benchmark interpretation: a per-family F1 gap is not interchangeable with a model deficit. Two of four families are bounded by gold quality, not model capacity; one (Prescription) is bounded by the model and mitigated by the deterministic lane on dev but not strictly dominated by it at full-200; one is a genuine recall target. Full provenance: docs/canon/workstreams/*_CANONICAL_LEDGER_CANON.md.

V. DISCUSSION

The two-task readout supports the value of modular, auditable extraction over a single monolithic score. Gan 2026 demonstrates that seizure-frequency extraction is a high-pressure temporal reasoning task: deterministic rules can reach high validation scores but show a pronounced holdout gap, while hybrid structured events provide the strongest frozen LLM-using aggregate evidence. ExECTv2 tests whether the same design discipline scales to broader clinical-fact recovery, where `clinical_headline` provides the headline surface and strict benchmark/CUI outputs remain diagnostic.

The most important reporting constraint is the separation of reliability from component impact. Reliability scorecards say whether a fixed architecture is grounded, calibrated, robust, consistent, and replayable under the declared surface; they do not prove that an individual component caused the observed performance. That causal question is reserved for the ablations and same-input replays of the component-impact analysis, bounded to the declared scorer, split, and inspection policy.

A. Limitations

Gan 2026 locked-test evidence is aggregate-only and does not authorize row-level test debugging or post-test tuning. ExECTv2 full-200 evidence is also aggregate-only under the current reliability and component-impact protocols. The ExECTv2 `clinical_headline` surface is intentionally not a strict benchmark win claim. Calibration and review-routing results are not deployment readiness claims, and Qwen repair v02 is model-family evidence rather than an operational promotion over GPT-4.1-mini or DeepSeek.

VI. CONCLUSION

Across deep seizure-frequency extraction and broad epilepsy-letter phenotyping, the modular architecture enables performance reporting, reliability assessment, and component-impact analysis without conflating their evidence boundaries. The current results support aggregate, claim-bounded conclusions: hybrid structured events are the strongest frozen Gan 2026 close-off candidate, and ExECTv2 achieves strong de-duplicated clinical-headline recovery with separate reliability and component-impact evidence.

APPENDIX A

ARTIFACT PROVENANCE AND CLAIM BOUNDARIES

APPENDIX B

SUPPORTING RESULT TABLES

The main text keeps four result tables: the Gan 2026 frozen aggregate audit, the ExECTv2 architecture comparison, the

ExECTv2 reliability scorecard, and the ExECTv2 component-impact table. The remaining four tables are supporting evidence: validation development context, late-cycle selector diagnostics, and same-core model-family robustness.

ACKNOWLEDGMENT

Acknowledgments omitted for draft.

REFERENCES

- [1] B. Fonferko-Shadrach, A. S. Lacey, A. Roberts, S. Thompson, D. V. Ford, R. A. Lyons, M. I. Rees, and W. O. Pickrell, "Using natural language processing to extract structured epilepsy data from unstructured clinic letters: development and validation of the ExECT (extraction of epilepsy clinical text) system," *BMJ Open*, vol. 9, no. 4, e023232, 2019, doi: 10.1136/bmjopen-2018-023232.
- [2] B. Fonferko-Shadrach *et al.*, "Annotation of epilepsy clinic letters for natural language processing," *Journal of Biomedical Semantics*, vol. 15, no. 17, 2024, doi: 10.1186/s13326-024-00316-z.
- [3] K. Xie *et al.*, "Extracting seizure frequency from epilepsy clinic notes: a machine reading approach to natural language processing," *Journal of the American Medical Informatics Association*, vol. 29, no. 5, pp. 873–881, 2022, doi: 10.1093/jamia/ocac018.
- [4] R. Abeysinghe, S. Tao, S. D. Lhatoo, G.-Q. Zhang, and L. Cui, "Leveraging pretrained language models for seizure frequency extraction from epilepsy evaluation reports," *npj Digital Medicine*, 2025, doi: 10.1038/s41746-025-01592-4.
- [5] B. Holgate, J. Davies, S. Fang, J. Winston, J. Teo, and M. Richardson, "Fine-tuning LLMs to extract epilepsy seizure frequency data from health records," in *Proc. 24th Workshop on Biomedical Language Processing*, Vienna, Austria, 2025, pp. 44–55, doi: 10.18653/v1/2025.bionlp-1.5.
- [6] Y. Gan, S. H. Barlow, B. Holgate, J. Davies, J. T. Teo, J. S. Winston, and M. P. Richardson, "Reproducible synthetic clinical letters for seizure frequency information extraction," arXiv: 2603.11407, 2026.

TABLE VI
MANUSCRIPT RESULT PROVENANCE AND CLAIM BOUNDARIES

Result family	Manuscript use	Source artifact anchor	Boundary
Gan validation architecture comparison	Appendix development context for the three-way comparison.	docs/research/paper_manuscript_2026-06-26.md; Gan close-off synthesis artifacts listed there.	Validation development evidence; not a frozen holdout claim.
Gan frozen architecture audit	Main frozen test450 comparison table.	experiments/gan2026_test450_phase4_comparison_report_gpt41mini_2026-06-10.md	One-shot locked-test aggregate audit; no row-level test tuning.
Gan consensus/fresh Gate 4 audits	Appendix late-cycle selector diagnostics.	experiments/gan2026_consensus_fresh_agreement_selector_v0_9_frozen_gate4_constrained_aggregate_audit_2026-06-26.md; experiments/gan2026_consensus_fresh_agreement_selector_v0_9_frozen_gate4_exact_aggregate_audit_2026-06-26.md	Frozen aggregate evidence only; constrained row failed, exact-source row passed within its source set.
ExECTv2 architecture and lean frontier	Main clinical_headline architecture table.	docs/experiments/exectv2/reliability/exectv2_holistic_finding_assembly_v08_full200_currentcode_gpt41mini_20260624.md; docs/experiments/exectv2/reliability/exectv2_gpt41mini_simplification_frontier_2026-06-24.md	Full-200 aggregate architecture evidence; no full-200 row-level development claim.
ExECTv2 same-core model swap	Appendix model-family robustness tables.	docs/experiments/exectv2/reliability/exectv2_same_core_model_swap_full200_2026-06-25.md; docs/experiments/exectv2/reliability/exectv2_qwen_model_swap_repair_v02_full200_readout_2026-06-26.md	Same-core aggregate model-family evidence; Qwen repair v02 is not operational promotion over GPT-4.1-mini or DeepSeek.
ExECTv2 reliability scorecard	Main reliability table and trust-language paragraph.	docs/research/exectv2_reliability_component_evidence_paper_language_2026-06-25.md; docs/experiments/exectv2/reliability/exectv2_calibration_validation_audit_2026-06-25.md; docs/experiments/exectv2/reliability/exectv2_robustness_validation_audit_2026-06-25.md	Reliability evidence for a fixed architecture; not holdout validation or component-causal proof.
ExECTv2 component impact	Main component-impact table.	experiments/exectv2_component_off_replay_full200_20260626.md; docs/experiments/exectv2/reliability/exectv2_component_off_full200_predeclaration_2026-06-26.md	Replay-only component-impact evidence under the declared scorer and aggregate-only boundary.

TABLE VII
GAN 2026 THREE-WAY COMPARISON ON VALIDATION750 (GPT-4.1-MINI)

Architecture	Rendered	Purist (of rendered)	Pragmatic (of rendered)	Reading
deterministic_canonical_pipeline	741/750	673/741 (0.908)	681/741 (0.919)	High validation score; large holdout drop.
hybrid	597/750	526/597 (0.881)	545/597 (0.913)	Strong rendered accuracy; too many null/routed rows.
hybrid_structured_events	748/750	661/748 (0.884)	679/748 (0.908)	Best LLM-using validation coverage/accuracy balance.
llm_only_canonical_pipeline	750/750	582/750 (0.776)	614/750 (0.819)	Comparator; below hybrid structured events.

TABLE VIII
GAN 2026 CONSENSUS/FRESH V0.9 FROZEN GATE 4 AGGREGATE AUDITS (TEST450)

Audit	Source symmetry	Selected Purist	Net gain	Changed precision	Gate outcome
Constrained	constrained	348/450 (0.7733)	+19	0.5909	Failed; final-evaluation evidence only.
Exact-source	exact	359/450 (0.7978)	+16	0.6000	Passed; frozen exact selector holdout.

TABLE IX
SAME-CORE MODEL SWAP ON DEV140 CLINICAL_HEADLINE

Candidate/model	Overall	Diagnosis	SF	Prescription	Investigations	Operational caveat
DeepSeek chat	0.8596	0.8845	0.7658	0.8895	0.8966	One parse/schema failure.
GPT-4.1-mini	0.8396	0.8573	0.7645	0.8895	0.8347	Operationally clean.
Qwen 3.6 35B repair v02	0.8319	0.8473	0.7182	0.8895	0.8755	Passes repair gates.
Qwen 3.6 35B unrepaired	0.8018	0.8027	0.6919	0.8895	0.8354	Diagnostic only.

TABLE X
SAME-CORE MODEL SWAP ON FULL-200 AGGREGATE CLINICAL_HEADLINE

Candidate/model	Overall	Diagnosis	SF	Prescription	Investigations	Call/parse failures
DeepSeek chat	0.8566	0.8708	0.7602	0.8926	0.9091	0 / 1
GPT-4.1-mini	0.8356	0.8397	0.7525	0.8926	0.8563	0 / 0
Qwen 3.6 35B repair v02	0.8197	0.8307	0.7020	0.8926	0.8503	0 / 0